

Part II

QR Factorization and Least Squares

Lecture 6. Projectors

We now enter the second part of the book, whose theme is orthogonality. We begin with the fundamental tool of projection matrices, or projectors, both orthogonal and nonorthogonal.

Projectors

A *projector* is a square matrix P that satisfies

$$P^2 = P. \quad (6.1)$$


(Such a matrix is also said to be *idempotent*.) This definition includes both orthogonal projectors, to be discussed in a moment, and nonorthogonal ones. To avoid confusion one may use the term oblique projector in the nonorthogonal case.

The term projector might be thought of as arising from the notion that if one were to shine a light onto the subspace $\text{range}(P)$ from just the right direction, then Pv would be the shadow projected by the vector v . We shall carry this physical picture forward for a moment.

Observe that if $v \in \text{range}(P)$, then it lies exactly on its own shadow, and applying the projector results in v itself. Mathematically, we have $v = Px$ for some x and

$$Pv = P^2x = Px = v.$$

From what direction does the light shine when $v \neq Pv$? In general the answer depends on v , but for any particular v , it is easily deduced by drawing the

Figure 6.1. An oblique projection.

line from v to Pv , $Pv - v$ (Figure 6.1). Applying the projector to this vector gives a zero result:

$$P(Pv - v) = P^2v - Pv = 0.$$

This means that $Pv - v \in \text{null}(P)$. That is, the direction of the light may be different for different v , but it is always described by a vector in $\text{null}(P)$.

Complementary Projectors

If P is a projector, $I - P$ is also a projector, for it is also idempotent:

$$(I - P)^2 = I - 2P + P^2 = I - P.$$

The matrix $I - P$ is called the complementary projector to P .

Onto what space does $I - P$ project? Exactly the nullspace of P ! We know that $\text{range}(I - P) \supseteq \text{null}(P)$, because if $Pv = 0$, we have $(I - P)v = v$. Conversely, we know that $\text{range}(I - P) \subseteq \text{null}(P)$, because for any v , we have $(I - P)v = v - Pv \in \text{null}(P)$. Therefore, for any projector P ,

$$\text{range}(I - P) = \text{null}(P). \quad (6.2)$$

By writing $P = I - (I - P)$ we derive the complementary fact


$$\text{null}(I - P) = \text{range}(P). \quad (6.3)$$

We can also see that $\text{null}(I - P) \cap \text{null}(P) = \{0\}$: any vector v in both sets satisfies $v = v - Pv = (I - P)v = 0$. Another way of stating this fact is

$$\text{range}(P) \cap \text{null}(P) = \{0\}. \quad (6.4)$$

These computations show that a projector separates $\mathbb{C}^m$ into two spaces. Conversely, let S_1 and S_2 be two subspaces of $\mathbb{C}^m$ such that $S_1 \cap S_2 = \{0\}$

P : the vectors in $S_1 \xrightarrow{(\text{ker})} S_2 \xrightarrow{(\text{im})}$
 $I - P$: the vectors in $S_2 \xrightarrow{(\text{ker})} S_1 \xrightarrow{(\text{im})}$






Figure 6.2. An orthogonal projection.

and $S_1 + S_2 = \mathbb{C}^m$, where $S_1 + S_2$ denotes the span of S_1 and S_2 , that is, the set of vectors $s_1 + s_2$ with $s_1 \in S_1$ and $s_2 \in S_2$. (Such a pair are said to be *complementary subspaces*.) Then there is a projector P such that $\text{range}(P) = S_1$ and $\text{null}(P) = S_2$. We say that P is the projector *onto* S_1 *along* S_2 . This projector and its complement can be seen as the unique solution to the following problem:

Given v , find vectors $v_1 \in S_1$ and $v_2 \in S_2$ such that $v_1 + v_2 = v$.

The projection Pv gives v_1 , and the complementary projection $(I - P)v$ gives v_2 . These vectors are unique because all solutions must be of the form

$$(Pv + v_3) + ((I - P)v - v_3) = v,$$

where it is clear that v_3 must be in both S_1 and S_2 , i.e., $v_3 = 0$.

One context in which projectors and their complements arise is particularly familiar. Suppose an $m \times m$ matrix A has a complete set of eigenvectors $\{v_j\}$, as in (5.1), meaning that $\{v_j\}$ is a basis of $\mathbb{C}^m$. We are frequently concerned with problems associated with expansions of vectors in this basis. Given $x \in \mathbb{C}^m$, for example, what is the component of x in the direction of a particular eigenvector v ? The answer is Px , where P is a certain rank-one projector. Rather than give details here, however, we turn now to the special class of projectors that will be of primary interest to us in this book.

Orthogonal Projectors

An *orthogonal projector* (Figure 6.2) is one that projects onto a subspace S_1 along a space S_2 , where S_1 and S_2 are orthogonal. (Warning: orthogonal projectors are not orthogonal matrices!)

Def orthogonal projector if:
 $\ker(P) \perp \text{im}(P)$

There is also an algebraic definition: an orthogonal projector is any projector that is hermitian, satisfying $P^* = P$ as well as (6.1). Of course, we must establish that this definition is equivalent to the first.

Theorem 6.1. A projector P is orthogonal if and only if $P = P^*$.

Proof. If $P = P^*$, then the inner product between a vector $Px \in S_1$ and a vector $(I - P)y \in S_2$ is zero:

$$\langle Px, (I - P)y \rangle = 0$$

$$x^* P^* (I - P)y = x^* (P - P^2)y = 0.$$

Thus the projector is orthogonal, providing the proof in the “if” direction.

For “only if,” we can use the SVD. Suppose P projects onto S_1 along S_2 , where $S_1 \perp S_2$ and S_1 has dimension n . Then an SVD of P can be constructed as follows. Let $\{q_1, q_2, \dots, q_m\}$ be an orthonormal basis for $\mathbb{C}^m$, where $\{q_1, \dots, q_n\}$ is a basis for S_1 and $\{q_{n+1}, \dots, q_m\}$ is a basis for S_2 . For $j \leq n$, we have $Pq_j = q_j$, and for $j > n$, we have $Pq_j = 0$. Now let Q be the unitary matrix whose j th column is q_j . We then have

$$PQ = \begin{bmatrix} | & | & | & | & | & | \\ q_1 & \cdots & q_n & 0 & \cdots & \\ | & | & | & | & | & | \end{bmatrix},$$

so that

$$Q^* P Q = \begin{bmatrix} 1 & & & & & \\ & \ddots & & & & \\ & & 1 & & & \\ & & & 0 & & \\ & & & & \ddots & \end{bmatrix} = \Sigma,$$

a diagonal matrix with ones in the first n entries and zeros everywhere else. Thus we have constructed a singular value decomposition of P :

$$P = Q \Sigma Q^*. \quad (6.5)$$

(Note that this is also an eigenvalue decomposition (5.1).) From here we see that P is hermitian, since $P^* = (Q \Sigma Q^*)^* = Q \Sigma^* Q^* = Q \Sigma Q^* = P$. $\square$

project u onto $\text{span}(v)$, $\|v\|_2 = 1$



length of proj: $l = v^* u (= u^* v)$

(since $|l| = \| \langle u, v \rangle v \| = \| \langle u, v \rangle \|$)

$$\text{so } \text{proj}_v(u) = l v = \begin{bmatrix} v \end{bmatrix} [v^* u] = v v^* u$$

Take outer product $P = v v^*$

$\Rightarrow P$ orthogonally project (arbitrary) u onto $\text{range}(v)$
if $v_1 \perp v_2$ unit

$$\Rightarrow \text{proj}_{v_1, v_2}(u) = v_1 v_1^* u + v_2 v_2^* u$$

$$\text{let } P = \sum_{i=1}^n v_i v_i^* \Rightarrow \boxed{\text{proj}_{v_1, \dots, v_n} u = P u}$$

Projection with an Orthonormal Basis

Since an orthogonal projector has some singular values equal to zero (except in the trivial case $P = I$), it is natural to drop the silent columns of Q in (6.5) and use the reduced rather than the full SVD. We obtain the marvelously simple expression

$$P = \hat{Q}\hat{Q}^*, \quad (6.6)$$

where the columns of $\hat{Q}$ are orthonormal.

In (6.6), the matrix $\hat{Q}$ need not come from an SVD. Let $\{q_1, \dots, q_n\}$ be any set of n orthonormal vectors in $\mathbb{C}^m$, and let $\hat{Q}$ be the corresponding $m \times n$ matrix. From (2.7) we know that

$$v = r + \sum_{i=1}^n (q_i q_i^*) v$$

r: remainder (distance from range(Q))

represents a decomposition of a vector $v \in \mathbb{C}^m$ into a component in the column space of $\hat{Q}$ plus a component in the orthogonal space. Thus the map

$$v \mapsto \sum_{i=1}^n (q_i q_i^*) v \quad (6.7)$$

is an orthogonal projector onto $\text{range}(\hat{Q})$, and in matrix form, it may be written $y = \hat{Q}\hat{Q}^*v$:

$$\begin{array}{c} \boxed{} \\ y \end{array} = \begin{array}{c} \boxed{} \\ \hat{Q} \end{array} \begin{array}{c} \boxed{} \\ \hat{Q}^* \end{array} \begin{array}{c} \boxed{} \\ v \end{array}.$$

Thus any product $\hat{Q}\hat{Q}^*$ is always a projector onto the column space of $\hat{Q}$, regardless of how $\hat{Q}$ was obtained, as long as its columns are orthonormal. Perhaps $\hat{Q}$ was obtained by dropping some columns and rows from a full factorization $v = QQ^*v$ of the identity,

$$\begin{array}{c} \boxed{} \\ v \end{array} = \begin{array}{c} \boxed{} \\ Q \end{array} \begin{array}{c} \boxed{} \\ Q^* \end{array} \begin{array}{c} \boxed{} \\ v \end{array},$$

and perhaps it was not.

The complement of an orthogonal projector is also an orthogonal projector (proof: $I - \hat{Q}\hat{Q}^*$ is hermitian). The complement projects onto the space orthogonal to $\text{range}(\hat{Q})$.

pf of: complement of an orthogonal proj is Hermitian
 $(I - QQ^*)^* = I^* - QQ^*$

($\Rightarrow$ thus, complement of orthogonal P is also orthogonal)

An important special case of orthogonal projectors is the rank-one orthogonal projector that isolates the component in a single direction q , which can be written

$$P_q = qq^*. \quad (\text{rank } 1) \quad (6.8)$$

These are the pieces from which higher-rank projectors can be made, as in (6.7). Their complements are the rank $m - 1$ orthogonal projectors that eliminate the component in the direction of q :

$$P_{\perp q} = I - qq^*. \quad (\text{rank } m-1) \quad (6.9)$$

Equations (6.8) and (6.9) assume that q is a unit vector. For arbitrary nonzero vectors a , the analogous formulas are

$$P_a = \frac{aa^*}{a^*a}, \quad (6.10)$$

$$P_{\perp a} = I - \frac{aa^*}{a^*a}. \quad (6.11)$$

Projection with an Arbitrary Basis

An orthogonal projector onto a subspace of $\mathbb{C}^m$ can also be constructed beginning with an arbitrary basis, not necessarily orthogonal. Suppose that the subspace is spanned by the linearly independent vectors $\{a_1, \dots, a_n\}$, and let A be the $m \times n$ matrix whose j th column is a_j .

In passing from v to its orthogonal projection $y \in \text{range}(A)$, the difference $y - v$ must be orthogonal to $\text{range}(A)$. This is equivalent to the statement that y must satisfy $a_j^*(y - v) = 0$ for every j . Since $y \in \text{range}(A)$, we can set $y = Ax$ and write this condition as $a_j^*(Ax - v) = 0$ for each j , or equivalently, $A^*(Ax - v) = 0$ or $A^*Ax = A^*v$. It is easily shown that since A has full rank, A^*A is nonsingular (Exercise 6.3). Therefore


$$x = (A^*A)^{-1}A^*v. \quad (6.12)$$

Finally, the projection of v , $y = Ax$, is $y = A(A^*A)^{-1}A^*v$. Thus the orthogonal projector onto $\text{range}(A)$ can be expressed by the formula

$$P = A(A^*A)^{-1}A^*. \quad (6.13)$$

Note that this is a multidimensional generalization of (6.10). In the orthonormal case $A = \hat{Q}$, the term in parentheses collapses to the identity and we recover (6.6).

want: $(y-v) \perp \text{im}(A)$
 i.e. $(y-v) \perp a_j \forall j$
 i.e. $a_j^*(y-v) = 0 \forall j$
 i.e. $A^*(y-v) = 0$


 (APv)
 suppose $y = Ax \Rightarrow$ sol. $x = (A^*A)^{-1}A^*v$
 (since $y \in \text{im}(A)$) $\Rightarrow P = A(A^*A)^{-1}A^*$

$[a_1 \dots a_n]$
 $:= A$

Exercises

6.1. If P is an orthogonal projector, then $I - 2P$ is unitary. Prove this algebraically, and give a geometric interpretation.

6.2. Let E be the $m \times m$ matrix that extracts the “even part” of an m -vector: $Ex = (x + Fx)/2$, where F is the $m \times m$ matrix that flips $(x_1, \dots, x_m)^*$ to $(x_m, \dots, x_1)^*$. Is E an orthogonal projector, an oblique projector, or not a projector at all? What are its entries?

6.3. Given $A \in \mathbb{C}^{m \times n}$ with $m \geq n$, show that A^*A is nonsingular if and only if A has full rank.

6.4. Consider the matrices

$$A = \begin{bmatrix} 1 & 0 \\ 0 & 1 \\ 1 & 0 \end{bmatrix}, \quad B = \begin{bmatrix} 1 & 2 \\ 0 & 1 \\ 1 & 0 \end{bmatrix}.$$

Answer the following questions by hand calculation.

(a) What is the orthogonal projector P onto $\text{range}(A)$, and what is the image under P of the vector $(1, 2, 3)^*$?

(b) Same questions for B .

6.5. Let $P \in \mathbb{C}^{m \times m}$ be a nonzero projector. Show that $\|P\|_2 \geq 1$, with equality if and only if P is an orthogonal projector.

Lecture 7. QR Factorization

reduced QR: for practice

full QR: for mathematical generality

One algorithmic idea in numerical linear algebra is more important than all the others: QR factorization.

Reduced QR Factorization

For many applications, we find ourselves interested in the column spaces of a matrix A . Note the plural: these are the *successive* spaces spanned by the columns $a_1, a_2, \dots$ of A :

$$\langle a_1 \rangle \subseteq \langle a_1, a_2 \rangle \subseteq \langle a_1, a_2, a_3 \rangle \subseteq \dots$$

Here, as in Lecture 5 and throughout the book, the notation $\langle \dots \rangle$ indicates the subspace spanned by whatever vectors are included in the brackets. Thus $\langle a_1 \rangle$ is the one-dimensional space spanned by a_1 , $\langle a_1, a_2 \rangle$ is the two-dimensional space spanned by a_1 and a_2 , and so on. The idea of QR factorization is the construction of a sequence of orthonormal vectors $q_1, q_2, \dots$ that span these successive spaces.

To be precise, assume for the moment that $A \in \mathbb{C}^{m \times n}$ ($m \geq n$) has full rank n . We want the sequence $q_1, q_2, \dots$ to have the property

$$\langle q_1, q_2, \dots, q_j \rangle = \langle a_1, a_2, \dots, a_j \rangle, \quad j = 1, \dots, n. \quad (7.1)$$

From the observations of Lecture 1, it is not hard to see that this amounts to

the condition

$$\begin{bmatrix} | & | & & | \\ a_1 & a_2 & \cdots & a_n \\ | & | & & | \end{bmatrix} = \begin{bmatrix} | & | & & | \\ q_1 & q_2 & \cdots & q_n \\ | & | & & | \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & \cdots & r_{1n} \\ & r_{22} & & \vdots \\ & & \ddots & \\ & & & r_{nn} \end{bmatrix}, \quad (7.2)$$

rank = n

where the diagonal entries r_{kk} are nonzero—for if (7.2) holds, then $a_1, \dots, a_k$ can be expressed as linear combinations of $q_1, \dots, q_k$, and the invertibility of the upper-left $k \times k$ block of the triangular matrix implies that, conversely, $q_1, \dots, q_k$ can be expressed as linear combinations of $a_1, \dots, a_k$. Written out, these equations take the form

$$\begin{aligned} a_1 &= r_{11}q_1, \\ a_2 &= r_{12}q_1 + r_{22}q_2, \\ a_3 &= r_{13}q_1 + r_{23}q_2 + r_{33}q_3, \\ &\vdots \\ a_n &= r_{1n}q_1 + r_{2n}q_2 + \cdots + r_{nn}q_n. \end{aligned} \quad (7.3)$$


As a matrix formula, we have

$$A = \hat{Q}\hat{R}, \quad (7.4)$$

where $\hat{Q}$ is $m \times n$ with orthonormal columns and $\hat{R}$ is $n \times n$ and upper-triangular. Such a factorization is called a *reduced QR factorization* of A .

Full QR Factorization

A *full QR factorization* of $A \in \mathbb{C}^{m \times n}$ ($m \geq n$) goes further, appending an additional $m - n$ orthonormal columns to $\hat{Q}$ so that it becomes an $m \times m$ unitary matrix Q . This is analogous to the passage from the reduced to the full SVD described in Lecture 4. In the process, rows of zeros are appended to $\hat{R}$ so that it becomes an $m \times n$ matrix R , still upper-triangular. The relationship between the full and reduced QR factorizations is as follows.



From (7.5) it is evident that an appropriate definition for the coefficients r_{ij} in the numerators of (7.6) is

$$r_{ij} = q_i^* a_j \quad (i \neq j). \quad (7.7)$$

The coefficients r_{jj} in the denominators are chosen for normalization:

$$|r_{jj}| = \left\| a_j - \sum_{i=1}^{j-1} r_{ij} q_i \right\|_2. \quad (7.8)$$

Note that the sign of r_{jj} is not determined. Arbitrarily, we may choose $r_{jj} > 0$, in which case we shall finish with a factorization $A = \hat{Q}\hat{R}$ in which $\hat{R}$ has positive entries along the diagonal.

The algorithm embodied in (7.6)–(7.8) is the Gram–Schmidt iteration. Mathematically, it offers a simple route to understanding and proving various properties of QR factorizations. Numerically, it turns out to be unstable because of rounding errors on a computer. To emphasize the instability, numerical analysts refer to this as the *classical Gram–Schmidt iteration*, as opposed to the *modified Gram–Schmidt iteration*, discussed in the next lecture.

Algorithm 7.1. Classical Gram–Schmidt (unstable)

for $j = 1$ **to** n

$v_j = a_j$

for $i = 1$ **to** $j - 1$

$r_{ij} = q_i^* a_j$

$v_j = v_j - r_{ij} q_i$

$r_{jj} = \|v_j\|_2$

$q_j = v_j / r_{jj}$

$$r_{ij} q_i = \text{proj}_{q_i} a_j$$

after $i = j-1$, $v_j = v_j - \text{span}(q_1, \dots, q_{j-1}) v_j$

Existence and Uniqueness

All matrices have QR factorizations, and under suitable restrictions, they are unique. We state first the existence result.

Theorem 7.1. Every $A \in \mathbb{C}^{m \times n}$ ($m \geq n$) has a full QR factorization, hence also a reduced QR factorization.

Proof. Suppose first that A has full rank and that we want just a reduced QR factorization. In this case, a proof of existence is provided by the Gram–Schmidt algorithm itself. By construction, this process generates orthonormal columns of $\hat{Q}$ and entries of $\hat{R}$ such that (7.4) holds. Failure can occur only if at some step, v_j is zero and thus cannot be normalized to produce q_j .

PF by
Gram-Schmidt
alg itself

However, this would imply $a_j \in \langle q_1, \dots, q_{j-1} \rangle = \langle a_1, \dots, a_{j-1} \rangle$, contradicting the assumption that A has full rank.

Now suppose that A does not have full rank. Then at one or more steps j , we shall find that (7.5) gives $v_j = 0$ as just mentioned. At this moment, we simply pick q_j arbitrarily to be any normalized vector orthogonal to $\langle q_1, \dots, q_{j-1} \rangle$, and then continue the Gram–Schmidt process.

Finally, the full, rather than reduced, QR factorization of an $m \times n$ matrix with $m > n$ can be constructed by introducing arbitrary orthonormal vectors in the same fashion. We follow the Gram–Schmidt process through step n , then continue on an additional $m - n$ steps, introducing vectors q_j at each step.

The issues discussed in the last two paragraphs came up already in Lecture 4, in our discussion of the SVD. $\square$

We turn now to uniqueness. Suppose $A = \hat{Q}\hat{R}$ is a reduced QR factorization. If the i th column of $\hat{Q}$ is multiplied by z and the i th row of $\hat{R}$ is multiplied by z^{-1} for some scalar z with $|z| = 1$, we obtain another QR factorization of A . The next theorem asserts that if A has full rank, this is the only way to obtain distinct reduced QR factorizations.

Theorem 7.2. *Each $A \in \mathbb{C}^{m \times n}$ ($m \geq n$) of full rank has a unique reduced QR factorization $A = \hat{Q}\hat{R}$ with $r_{jj} > 0$.*

Proof. Again, the proof is provided by the Gram–Schmidt iteration. From (7.4), the orthonormality of the columns of $\hat{Q}$, and the upper-triangularity of $\hat{R}$, it follows that any reduced QR factorization of A must satisfy (7.6)–(7.8). By the assumption of full rank, the denominators (7.8) of (7.6) are nonzero, and thus at each successive step j , these formulas determine r_{ij} and q_j fully, except in one place: the sign of r_{jj} , not specified in (7.8). Once this is fixed by the condition $r_{jj} > 0$, as in Algorithm 7.1, the factorization is completely determined. $\square$

When Vectors Become Continuous Functions

The QR factorization has an analogue for orthonormal expansions of functions rather than vectors.

Suppose we replace $\mathbb{C}^m$ by $L^2[-1, 1]$, a vector space of complex-valued functions on $[-1, 1]$. We shall not introduce the properties of this space formally; suffice it to say that the inner product of f and g now takes the form

$$(f, g) = \int_{-1}^1 \overline{f(x)} g(x) dx. \quad (7.9)$$

Consider, for example, the following “matrix” whose “columns” are the monomials x^j :

$$A = \left[\begin{array}{c|c|c|c|c} 1 & x & x^2 & \cdots & x^{n-1} \end{array} \right]. \quad (7.10)$$

Each column is a function in $L^2[-1, 1]$, and thus, whereas A is discrete as usual in the horizontal direction, it is continuous in the vertical direction. It is a continuous analogue of the Vandermonde matrix (1.4) of Example 1.1.

The “continuous QR factorization” of A takes the form

$$A = QR = \left[\begin{array}{c|c|c|c|c} q_0(x) & q_1(x) & \cdots & q_{n-1}(x) \end{array} \right] \begin{bmatrix} r_{11} & r_{12} & \cdots & r_{1n} \\ & r_{22} & & \vdots \\ & & \ddots & r_{nn} \end{bmatrix},$$

$[-1, 1] \times \mathbb{R} \quad \mathbb{R} \times \mathbb{R}$

where the columns of Q are functions of x , orthonormal with respect to the inner product (7.9):

$$\int_{-1}^1 \overline{q_i(x)} q_j(x) dx = \delta_{ij} = \begin{cases} 1 & \text{if } i = j, \\ 0 & \text{if } i \neq j. \end{cases}$$

From the Gram–Schmidt construction we can see that q_j is a polynomial of degree j . These polynomials are scalar multiples of what are known as the Legendre polynomials, P_j , which are conventionally normalized so that $P_j(1) = 1$. The first few P_j are

$$P_0(x) = 1, \quad P_1(x) = x, \quad P_2(x) = \frac{3}{2}x^2 - \frac{1}{2}, \quad P_3(x) = \frac{5}{2}x^3 - \frac{3}{2}x; \quad (7.11)$$

see Figure 7.1. Like the monomials $1, x, x^2, \dots$, this sequence of polynomials spans the spaces of polynomials of successively higher degree. However, $P_0(x), P_1(x), P_2(x), \dots$ have the advantage that they are orthogonal, making them far better suited for certain computations. In fact, computations with such polynomials form the basis of spectral methods, one of the most powerful techniques for the numerical solution of partial differential equations.

What is the “projection matrix” $\hat{Q}\hat{Q}^*$ (6.6) associated with $\hat{Q}$? It is a “[$-1, 1$] $\times$ [$-1, 1$] matrix,” that is, an integral operator

$$f(\cdot) \mapsto \sum_{j=0}^{n-1} q_j(\cdot) \int_{-1}^1 \overline{q_j(x)} f(x) dx \quad (7.12)$$

mapping functions in $L^2[-1, 1]$ to functions in $L^2[-1, 1]$.




Figure 7.1. The first four Legendre polynomials (7.11). Apart from scale factors, these can be interpreted as the columns of $\hat{Q}$ in a reduced QR factorization of the “ $[-1, 1] \times 4$ matrix” $[1, x, x^2, x^3]$.

Solution of $Ax = b$ by QR Factorization

In closing this lecture we return for a moment to discrete, finite matrices. Suppose we wish to solve $Ax = b$ for x , where $A \in \mathbb{C}^{m \times m}$ is nonsingular. If $A = QR$ is a QR factorization, then we can write $QRx = b$, or

$$\underline{Rx = Q^*b.} \quad (7.13)$$

The right-hand side of this equation is easy to compute, if Q is known, and the system of linear equations implicit in the left-hand side is also easy to solve because it is triangular. This suggests the following method for computing the solution to $Ax = b$:

1. Compute a QR factorization $A = QR$.
2. Compute $y = Q^*b$.
3. Solve $Rx = y$ for x .

(at least faster than

In later lectures we shall present algorithms for each of these steps.

The combination 1–3 is an excellent method for solving linear systems of equations; in Lecture 16, we shall prove this. However, it is not the standard method for such problems. Gaussian elimination is the algorithm generally used in practice, since it requires only half as many numerical operations.

Exercises

7.1. Consider again the matrices A and B of Exercise 6.4.

(a) Using any method you like, determine (on paper) a reduced QR factorization $A = \hat{Q}\hat{R}$ and a full QR factorization $A = QR$.

(b) Again using any method you like, determine reduced and full QR factorizations $B = \hat{Q}\hat{R}$ and $B = QR$.

7.2. Let A be a matrix with the property that columns 1, 3, 5, 7, ... are orthogonal to columns 2, 4, 6, 8, In a reduced QR factorization $A = \hat{Q}\hat{R}$, what special structure does $\hat{R}$ possess? You may assume that A has full rank.

7.3. Let A be an $m \times m$ matrix, and let a_j be its j th column. Give an algebraic proof of *Hadamard's inequality*:

$$|\det A| \leq \prod_{j=1}^m \|a_j\|_2.$$

Also give a geometric interpretation of this result, making use of the fact that the determinant equals the volume of a parallelepiped.

7.4. Let $x^{(1)}$, $y^{(1)}$, $x^{(2)}$, and $y^{(2)}$ be nonzero vectors in $\mathbb{R}^3$ with the property that $x^{(1)}$ and $y^{(1)}$ are linearly independent and so are $x^{(2)}$ and $y^{(2)}$. Consider the two planes in $\mathbb{R}^3$,

$$P^{(1)} = \langle x^{(1)}, y^{(1)} \rangle, \quad P^{(2)} = \langle x^{(2)}, y^{(2)} \rangle.$$

Suppose we wish to find a nonzero vector $v \in \mathbb{R}^3$ that lies in the intersection $P = P^{(1)} \cap P^{(2)}$. Devise a method for solving this problem by reducing it to the computation of QR factorizations of three 3×2 matrices.

7.5. Let A be an $m \times n$ matrix ($m \geq n$), and let $A = \hat{Q}\hat{R}$ be a reduced QR factorization.

(a) Show that A has rank n if and only if all the diagonal entries of $\hat{R}$ are nonzero.

(b) Suppose $\hat{R}$ has k nonzero diagonal entries for some k with $0 \leq k < n$. What does this imply about the rank of A ? Exactly k ? At least k ? At most k ? Give a precise answer, and prove it.

Lecture 8. Gram–Schmidt Orthogonalization

The Gram–Schmidt iteration is the basis of one of the two principal numerical algorithms for computing QR factorizations. It is a process of “triangular orthogonalization,” making the columns of a matrix orthonormal via a sequence of matrix operations that can be interpreted as multiplication on the right by upper-triangular matrices.

Gram–Schmidt Projections

In the last lecture we presented the Gram–Schmidt iteration in its classical form. To begin this lecture, we describe the same algorithm again in another way, using orthogonal projectors.

Let $A \in \mathbb{C}^{m \times n}$, $m \geq n$, be a matrix of full rank with columns $\{a_j\}$. Before, we expressed the Gram–Schmidt iteration by the formulas (7.6)–(7.8). Consider now the sequence of formulas


$$q_1 = \frac{P_1 a_1}{\|P_1 a_1\|}, \quad q_2 = \frac{P_2 a_2}{\|P_2 a_2\|}, \quad \dots, \quad q_n = \frac{P_n a_n}{\|P_n a_n\|}. \quad (8.1)$$

In these formulas, each P_j denotes an orthogonal projector. Specifically, P_j is the $m \times m$ matrix of rank $m - (j - 1)$ that projects $\mathbb{C}^m$ orthogonally onto the space orthogonal to $\langle q_1, \dots, q_{j-1} \rangle$. (In the case $j = 1$, this prescription reduces to the identity: $P_1 = I$.) Now, observe that q_j as defined by (8.1) is

orthogonal to $q_1, \dots, q_{j-1}$, lies in the space $\langle a_1, \dots, a_j \rangle$, and has norm 1. Thus we see that (8.1) is equivalent to (7.6)–(7.8) and hence to Algorithm 7.1.

The projector P_j can be represented explicitly. Let $\hat{Q}_{j-1}$ denote the $m \times (j-1)$ matrix containing the first $j-1$ columns of $\hat{Q}$,

$$\hat{Q}_{j-1} = \begin{bmatrix} | & | & & | \\ q_1 & q_2 & \cdots & q_{j-1} \\ | & | & & | \end{bmatrix}.$$



Then P_j is given by

$$P_j = I - \underbrace{\hat{Q}_{j-1} \hat{Q}_{j-1}^*}_{\text{proj}_{\text{span}(q_1, \dots, q_{j-1})}}. \quad (8.3)$$

By now, the reader may be familiar enough with our notation and with orthogonality ideas to see at a glance that (8.3) represents the operator applied to a_j in (7.5).

Modified Gram-Schmidt Algorithm

proj_{span^\perp(q_1, \dots, q_{j-1})}

In practice, the Gram-Schmidt formulas are not applied as we have indicated in Algorithm 7.1 and in (8.1), for this sequence of calculations turns out to be numerically unstable. Fortunately, there is a simple modification that improves matters. We have not discussed numerical stability yet; this will come in the next lecture and then systematically beginning in Lecture 14. For the moment, it is enough to know that a stable algorithm is one that is not too sensitive to the effects of rounding errors on a computer.

For each value of j , Algorithm 7.1 computes a single orthogonal projection of rank $m - (j - 1)$,

$$v_j = P_j a_j. \quad (8.4)$$

In contrast, the modified Gram-Schmidt algorithm computes the same result by a sequence of $j - 1$ projections of rank $m - 1$. Recall from (6.9) that $P_{\perp q}$ denotes the rank $m - 1$ orthogonal projector onto the space orthogonal to a nonzero vector $q \in \mathbb{C}^m$. By the definition of P_j , it is not difficult to see that

$$P_j = P_{\perp q_{j-1}} \cdots P_{\perp q_2} P_{\perp q_1}, \quad (8.5)$$

again with $P_1 = I$. Thus an equivalent statement to (8.4) is

$$v_j = P_{\perp q_{j-1}} \cdots P_{\perp q_2} P_{\perp q_1} a_j. \quad (8.6)$$

The modified Gram-Schmidt algorithm is based on the use of (8.6) instead of (8.4).

Mathematically, (8.6) and (8.4) are equivalent. However, the sequences of arithmetic operations implied by these formulas are different. The modified algorithm calculates v_j by evaluating the following formulas in order:

$$\begin{aligned}
 v_j^{(1)} &= a_j, \\
 v_j^{(2)} &= P_{\perp q_1} v_j^{(1)} = v_j^{(1)} - q_1 q_1^* v_j^{(1)}, \\
 v_j^{(3)} &= P_{\perp q_2} v_j^{(2)} = v_j^{(2)} - q_2 q_2^* v_j^{(2)}, \\
 &\vdots \\
 v_j &= v_j^{(j)} = P_{\perp q_{j-1}} v_j^{(j-1)} = v_j^{(j-1)} - q_{j-1} q_{j-1}^* v_j^{(j-1)}.
 \end{aligned} \tag{8.7}$$

In finite precision computer arithmetic, we shall see that (8.7) introduces smaller errors than (8.4).

When the algorithm is implemented, the projector $P_{\perp q_i}$ can be conveniently applied to $v_j^{(i)}$ for each $j > i$ immediately after q_i is known. This is done in the description below.

Algorithm 8.1. Modified Gram–Schmidt

```

for i = 1 to n
    v_i = a_i
for i = 1 to n
    r_ii = ||v_i||
    q_i = v_i / r_ii
    for j = i + 1 to n
        r_ij = q_i^* v_j
        v_j = v_j - r_ij q_i

```

	v_1	v_2	$\dots$	v_n
1	a_1	a_2	$\dots$	a_n
2	a_1	$a_2 - \text{proj}_{q_1} a_2$	$\dots$	$a_n - \text{proj}_{q_1} a_n$
3	a_1	$a_2 - \text{proj}_{q_1} a_2 - \text{proj}_{q_2} (a_2 - \text{proj}_{q_1} a_2)$	$\dots$	$a_n - \text{proj}_{q_1} a_n - \text{proj}_{q_2} (a_n - \text{proj}_{q_1} a_n)$

In practice, it is common to let v_i overwrite a_i and q_i overwrite v_i in order to save storage.

The reader should compare Algorithms 7.1 and 8.1 until he or she is confident of their equivalence.

Operation Count

The Gram–Schmidt algorithm is the first algorithm we have presented in this book, and with any algorithm, it is important to assess its cost. To do so, throughout the book we follow the classical route and count the number of floating point operations—“*flops*”—that the algorithm requires. Each addition, subtraction, multiplication, division, or square root counts as one flop.

We make no distinction between real and complex arithmetic, although in practice on most computers there is a sizable difference.

In fact, there is much more to the cost of an algorithm than operation counts. On a single-processor computer, the execution time is affected by the movement of data between elements of the memory hierarchy and by competing jobs running on the same processor. On multiprocessor machines the situation becomes more complex, with communication between processors sometimes taking on an importance much greater than that of actual “computation.” With some regret, we shall ignore these important considerations, because this book is deliberately classical in style, focusing on algorithmic foundations.

For both variants of the Gram–Schmidt iteration, here is the classical result.

Theorem 8.1. *Algorithms 7.1 and 8.1 require $\sim 2mn^2$ flops to compute a QR factorization of an $m \times n$ matrix.*

Note that the theorem expresses only the leading term of the flop count. The symbol “ $\sim$ ” has its usual asymptotic meaning:

$$\lim_{m,n \rightarrow \infty} \frac{\text{number of flops}}{2mn^2} = 1.$$

In discussing operation counts for algorithms, it is standard to discard lower-order terms as we have done here, since they are usually of little significance unless m and n are small.

Theorem 8.1 can be established as follows. To be definite, consider the modified Gram–Schmidt algorithm, Algorithm 8.1. When m and n are large, the work is dominated by the operations in the innermost loop:

$$\begin{aligned} r_{ij} &= q_i^* v_j, \\ v_j &= v_j - r_{ij} q_i. \end{aligned}$$


The first line computes an inner product $q_i^* v_j$, requiring m multiplications and $m-1$ additions, and the second computes $v_j - r_{ij} q_i$, requiring m multiplications and m subtractions. The total work involved in a single inner iteration is consequently $\sim 4m$ flops, or 4 flops per column vector element. All together, the number of flops required by the algorithm is asymptotic to

$$\sum_{i=1}^n \sum_{j=i+1}^n 4m \sim \sum_{i=1}^n (i) 4m \sim 2mn^2. \quad (8.8)$$

Counting Operations Geometrically

Operation counts can always be determined algebraically as in (8.8), and this is the standard procedure in the numerical analysis literature. However, it is

also enlightening to take a different, geometrical route to the same conclusion. The argument goes like this. At the first step of the outer loop, Algorithm 8.1 operates on the whole matrix, subtracting a multiple of column 1 from the other columns. At the second step, it operates on a submatrix, subtracting a multiple of column 2 from columns 3, $\dots$, n . Continuing on in this way, at each step the column dimension shrinks by 1 until at the final step, only column n is modified. This process can be represented by the following diagram:



The $m \times n$ rectangle at the bottom corresponds to the first pass through the outer loop, the $m \times (n - 1)$ rectangle above it to the second pass, and so on.

To leading order as $m, n \rightarrow \infty$, then, the operation count for Gram-Schmidt orthogonalization is proportional to the volume of the figure above. The constant of proportionality is four flops, because as noted above, the two steps of the inner loop correspond to four operations at each matrix location. Now as $m, n \rightarrow \infty$, the figure converges to a right triangular prism, with volume $mn^2/2$. Multiplying by four flops per unit volume gives, again,

$$\text{Work for Gram-Schmidt orthogonalization: } \sim 2mn^2 \text{ flops.} \quad (8.9)$$

In this book we generally record operation counts in the format (8.9), without stating them as theorems. We often derive these results via figures like the one above, although algebraic derivations are also possible. One reason we do this is that a figure of this kind, besides being a route to an operation count, also serves as a reminder of the structure of an algorithm. For pictures of algorithms with different structures, see pp. 75 and 176.

Gram–Schmidt as Triangular Orthogonalization

Each outer step of the modified Gram–Schmidt algorithm can be interpreted as a right-multiplication by a square upper-triangular matrix. For example, beginning with A , the first iteration multiplies the first column a_1 by $1/r_{11}$ and then subtracts r_{1j} times the result from each of the remaining columns a_j . This is equivalent to right-multiplication by a matrix R_1 :

$$\left[\begin{array}{c|c|c|c} v_1 & v_2 & \cdots & v_n \end{array} \right] \begin{bmatrix} \frac{1}{r_{11}} & \frac{-r_{12}}{r_{11}} & \frac{-r_{13}}{r_{11}} & \cdots \\ & 1 & & \\ & & 1 & \\ & & & \ddots \end{bmatrix} = \left[\begin{array}{c|c|c|c} q_1 & v_2^{(2)} & \cdots & v_n^{(2)} \end{array} \right].$$

In general, step i of Algorithm 8.1 subtracts r_{ij}/r_{ii} times column i of the current A from columns $j > i$ and replaces column i by $1/r_{ii}$ times itself. This corresponds to multiplication by an upper-triangular matrix R_i :

$$R_2 = \begin{bmatrix} 1 & & & \\ & \frac{1}{r_{22}} & \frac{-r_{23}}{r_{22}} & \cdots \\ & & 1 & \\ & & & \ddots \end{bmatrix}, \quad R_3 = \begin{bmatrix} 1 & & & \\ & 1 & & \\ & & \frac{1}{r_{33}} & \cdots \\ & & & \ddots \end{bmatrix}, \dots$$

At the end of the iteration we have

$$A \underbrace{R_1 R_2 \cdots R_n}_{\hat{R}^{-1}} = \hat{Q}. \quad (8.10)$$

This formulation demonstrates that the Gram–Schmidt algorithm is a method of *triangular orthogonalization*. It applies triangular operations on the right of a matrix to reduce it to a matrix with orthonormal columns. Of course, in practice, we do not form the matrices R_i and multiply them together explicitly. The purpose of mentioning them is to give insight into the structure of the Gram–Schmidt algorithm. In Lecture 20 we shall see that it bears a close resemblance to the structure of Gaussian elimination.

Exercises

8.1. Let A be an $m \times n$ matrix. Determine the exact numbers of floating point additions, subtractions, multiplications, and divisions involved in computing the factorization $A = \hat{Q}\hat{R}$ by Algorithm 8.1.

8.2. Write a MATLAB function $[Q,R] = \text{mgs}(A)$ (see next lecture) that computes a reduced QR factorization $A = \hat{Q}\hat{R}$ of an $m \times n$ matrix A with $m \geq n$ using modified Gram–Schmidt orthogonalization. The output variables are a matrix $Q \in \mathbb{C}^{m \times n}$ with orthonormal columns and a triangular matrix $R \in \mathbb{C}^{n \times n}$.

8.3. Each upper-triangular matrix R_j of p. 61 can be interpreted as the product of a diagonal matrix and a unit upper-triangular matrix (i.e., an upper-triangular matrix with 1 on the diagonal). Explain exactly what these factors are, and which line of Algorithm 8.1 corresponds to each.

Lecture 10. Householder Triangularization

The other principal method for computing QR factorizations is Householder triangularization, which is numerically more stable than Gram–Schmidt orthogonalization, though it lacks the latter’s applicability as a basis for iterative methods. The Householder algorithm is a process of “orthogonal triangularization,” making a matrix triangular by a sequence of unitary matrix operations.

Householder and Gram–Schmidt

As we saw in Lecture 8, the Gram–Schmidt iteration applies a succession of elementary triangular matrices R_k on the right of A , so that the resulting matrix

$$A \underbrace{R_1 R_2 \cdots R_n}_{\hat{R}^{-1}} = \hat{Q}$$

has orthonormal columns. The product $\hat{R} = R_n^{-1} \cdots R_2^{-1} R_1^{-1}$ is upper-triangular too, and thus $A = \hat{Q} \hat{R}$ is a reduced QR factorization of A .

In contrast, the Householder method applies a succession of elementary unitary matrices Q_k on the left of A , so that the resulting matrix

$$\underbrace{Q_n \cdots Q_2 Q_1}_{Q^*} A = R$$

is upper-triangular. The product $Q = Q_1^* Q_2^* \cdots Q_n^*$ is unitary too, and therefore $A = QR$ is a full QR factorization of A .

The two methods can thus be summarized as follows:

Gram–Schmidt: triangular orthogonalization,
Householder: orthogonal triangularization.

Triangularizing by Introducing Zeros

At the heart of the Householder method is an idea originally proposed by Alston Householder in 1958. This is an ingenious way of designing the unitary matrices Q_k so that $Q_n \cdots Q_2 Q_1 A$ is upper-triangular.

The matrix Q_k is chosen to introduce zeros below the diagonal in the k th column while preserving all the zeros previously introduced. For example, in the 5×3 case, three operations Q_k are applied, as follows. In these matrices, the symbol $\times$ represents an entry that is not necessarily zero, and boldfacing indicates an entry that has just been changed. Blank entries are zero.

$$\begin{bmatrix} \times & \times & \times \\ \times & \times & \times \\ \times & \times & \times \\ \times & \times & \times \\ \times & \times & \times \end{bmatrix} \xrightarrow{Q_1} \begin{bmatrix} \times & \times & \times \\ \mathbf{0} & \times & \times \\ \mathbf{0} & \times & \times \\ \mathbf{0} & \times & \times \\ \mathbf{0} & \times & \times \end{bmatrix} \xrightarrow{Q_2} \begin{bmatrix} \times & \times & \times \\ & \times & \times \\ & \mathbf{0} & \times \\ & \mathbf{0} & \times \\ & \mathbf{0} & \times \end{bmatrix} \xrightarrow{Q_3} \begin{bmatrix} \times & \times & \times \\ & \times & \times \\ & & \times \\ & & \mathbf{0} \\ & & \mathbf{0} \end{bmatrix} \quad (10.1)$$

$A \qquad Q_1 A \qquad Q_2 Q_1 A \qquad Q_3 Q_2 Q_1 A$

First, Q_1 operates on rows 1, ..., 5, introducing zeros in positions (2, 1), (3, 1), (4, 1), and (5, 1). Next, Q_2 operates on rows 2, ..., 5, introducing zeros in positions (3, 2), (4, 2), and (5, 2) but not destroying the zeros introduced by Q_1 . Finally, Q_3 operates on rows 3, ..., 5, introducing zeros in positions (4, 3) and (5, 3) without destroying any of the zeros introduced earlier.

In general, Q_k operates on rows $k, \dots, m$. At the beginning of step k , there is a block of zeros in the first $k - 1$ columns of these rows. The application of Q_k forms linear combinations of these rows, and the linear combinations of the zero entries remain zero. After n steps, all the entries below the diagonal have been eliminated and $Q_n \cdots Q_2 Q_1 A = R$ is upper-triangular.

Householder Reflectors

How can we construct unitary matrices Q_k to introduce zeros as indicated in (10.1)? The standard approach is as follows. Each Q_k is chosen to be a unitary matrix of the form

$$Q_k = \begin{bmatrix} I & 0 \\ 0 & F \end{bmatrix}, \quad (10.2)$$

where I is the $(k - 1) \times (k - 1)$ identity and F is an $(m - k + 1) \times (m - k + 1)$ unitary matrix. Multiplication by F must introduce zeros into the

Q_k : 令
第 k 列
的 k 元素
以下清零
即 $A_{k:m, k:n}$
submatrix
的第一列元素 - pivot

保留前 $k-1$ 行和 $k-1$ 列的 entries
作用 $A_{k:m, k:n}$




Figure 10.1. A Householder reflection.

k th column. The Householder algorithm chooses F to be a particular matrix called a *Householder reflector*.


Suppose, at the beginning of step k , the entries $k, \dots, m$ of the k th column are given by the vector $x \in \mathbb{C}^{m-k+1}$. To introduce the correct zeros into the k th column, the Householder reflector F should effect the following map:

$$x = \begin{bmatrix} \times \\ \times \\ \times \\ \vdots \\ \times \end{bmatrix} \xrightarrow{F} Fx = \begin{bmatrix} \|x\| \\ 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix} = \|x\|e_1. \quad (10.3)$$

(We shall modify this idea by a $\pm$ sign in a moment.) The idea for accomplishing this is indicated in Figure 10.1. The reflector F will reflect the space $\mathbb{C}^{m-k+1}$ across the hyperplane H orthogonal to $v = \|x\|e_1 - x$. A *hyperplane* is the higher-dimensional generalization of a two-dimensional plane in three-space—a three-dimensional subspace of a four-dimensional space, a four-dimensional subspace of a five-dimensional space, and so on. In general, a hyperplane can be characterized as the set of points orthogonal to a fixed nonzero vector. In Figure 10.1, that vector is $v = \|x\|e_1 - x$, and one can think of the dashed line as a depiction of H viewed “edge on.”

When the reflector is applied, every point on one side of the hyperplane H is mapped to its mirror image on the other side. In particular, x is mapped to $\|x\|e_1$. The formula for this reflection can be derived as follows. In (6.11) we have seen that for any $y \in \mathbb{C}^m$, the vector

$$Py = \left(I - \frac{vv^*}{v^*v}\right)y = y - v\left(\frac{v^*y}{v^*v}\right)$$






Figure 10.2. Two possible reflections. For numerical stability, it is important to choose the one that moves x the larger distance.

is the orthogonal projection of y onto the space H . To reflect y across H , we must not stop at this point; we must go exactly twice as far in the same direction. The reflection Fy should therefore be

$$Fy = \left(I - 2 \frac{vv^*}{v^*v} \right) y = y - 2v \left(\frac{v^*y}{v^*v} \right).$$

Hence the matrix F is

$$F = I - 2 \frac{vv^*}{v^*v}. \quad (10.4)$$

Note that the projector P (rank $m-1$) and the reflector F (full rank, unitary) differ only in the presence of a factor of 2.

The Better of Two Reflectors

In (10.3) and in Figure 10.1 we have simplified matters, for in fact, there are many Householder reflections that will introduce the zeros needed. The vector x can be reflected to $z\|x\|e_1$, where z is any scalar with $|z| = 1$. In the complex case, there is a circle of possible reflections, and even in the real case, there are two alternatives, represented by reflections across two different hyperplanes, H^+ and H^- , as illustrated in Figure 10.2.

Mathematically, either choice of sign is satisfactory. However, this is a case where the goal of numerical stability—insensitivity to rounding errors—dictates that one choice should be taken rather than the other. For numerical stability, it is desirable to reflect x to the vector $z\|x\|e_1$ that is not too close to x itself. To achieve this, we can choose $z = -\text{sign}(x_1)$, where x_1 denotes the first component of x , so that the reflection vector becomes $v = -\text{sign}(x_1)\|x\|e_1 - x$,

or, upon clearing the factors -1 ,

$$v = \text{sign}(x_1)\|x\|e_1 + x. \quad (10.5)$$

To make this a complete prescription, we may arbitrarily impose the convention that $\text{sign}(x_1) = 1$ if $x_1 = 0$.

It is not hard to see why the choice of sign makes a difference for stability. Suppose that in Figure 10.2, the angle between H^+ and the e_1 axis is very small. Then the vector $v = \|x\|e_1 - x$ is much smaller than x or $\|x\|e_1$. Thus the calculation of v represents a subtraction of nearby quantities and will tend to suffer from cancellation errors. If we pick the sign as in (10.5), we avoid such effects by ensuring that $\|v\|$ is never smaller than $\|x\|$.

The Algorithm

We now formulate the whole Householder algorithm. To do this, it will be helpful to utilize a new (MATLAB-style) notation. If A is a matrix, we define $A_{i:i', j:j'}$ to be the $(i' - i + 1) \times (j' - j + 1)$ submatrix of A with upper-left corner a_{ij} and lower-right corner $a_{i', j'}$. In the special case where the submatrix reduces to a subvector of a single row or column, we write $A_{i, j:j'}$ or $A_{i:i', j}$, respectively.

The following algorithm computes the factor R of a QR factorization of an $m \times n$ matrix A with $m \geq n$, leaving the result in place of A . Along the way, n reflection vectors $v_1, \dots, v_n$ are stored for later use.

Algorithm 10.1. Householder QR Factorization

for $k = 1$ **to** n

$$x = A_{k:m, k}$$

$$v_k = \text{sign}(x_1)\|x\|_2 e_1 + x$$

$$v_k = v_k / \|v_k\|_2$$

$$A_{k:m, k:n} = A_{k:m, k:n} - 2v_k(v_k^* A_{k:m, k:n})$$

Applying or Forming Q

Upon the completion of Algorithm 10.1, A has been reduced to upper-triangular form; this is the matrix R in the QR factorization $A = QR$. The unitary matrix Q has not, however, been constructed, nor has its n -column submatrix $\hat{Q}$ corresponding to a reduced QR factorization. There is a reason for this. Constructing Q or $\hat{Q}$ takes additional work, and in many applications, we can avoid this by working directly with the formula

$$Q^* = Q_n \cdots Q_2 Q_1 \quad (10.6)$$

or its conjugate

$$Q = Q_1 Q_2 \cdots Q_n. \quad (10.7)$$

(No asterisks have been forgotten here; recall that each Q_j is hermitian.)

For example, in Lecture 7 we saw that a square system of equations $Ax = b$ can be solved via QR factorization of A . The only way in which Q was used in this process was in the computation of the product Q^*b . By (10.6), we can calculate Q^*b by a sequence of n operations applied to b , the same operations that were applied to A to make it triangular. The algorithm is as follows.

Algorithm 10.2. Implicit Calculation of a Product Q^*b

for $k = 1$ **to** n

$$b_{k:m} = b_{k:m} - 2v_k(v_k^* b_{k:m})$$

Similarly, the computation of a product Qx can be achieved by the same process executed in reverse order.

Algorithm 10.3. Implicit Calculation of a Product Qx

for $k = n$ **downto** 1

$$x_{k:m} = x_{k:m} - 2v_k(v_k^* x_{k:m})$$

The work involved in either of these algorithms is of order $O(mn)$, not $O(mn^2)$ as in Algorithm 10.1 (see below).

Sometimes, of course, one may wish to construct the matrix Q explicitly. This can be achieved in various ways. We can construct QI via Algorithm 10.3 by computing its columns $Qe_1, Qe_2, \dots, Qe_m$. Alternatively, we can construct Q^*I via Algorithm 10.2 and then conjugate the result. A variant of this idea is to conjugate each step rather than the final product, that is, to construct IQ by computing its rows $e_1^*Q, e_2^*Q, \dots, e_m^*Q$ as suggested by (10.7). Of these various ideas, the best is the first one, based on Algorithm 10.3. The reason is that it begins with operations involving Q_n, Q_{n-1} , and so on that modify only a small part of the vector they are applied to; if advantage is taken of this sparsity property, a speed-up is achieved.

If only $\hat{Q}$ rather than Q is needed, it is enough to compute the columns $Qe_1, Qe_2, \dots, Qe_n$.


Operation Count

The work involved in Algorithm 10.1 is dominated by the innermost loop,


$$A_{k:m,j} - 2v_k(v_k^* A_{k:m,j}). \quad (10.8)$$

If the vector length is $l = m - k + 1$, this calculation requires $4l - 1 \sim 4l$ scalar operations: l for the subtraction, l for the scalar multiplication, and $2l - 1$ for the dot product. This is ~ 4 flops for each entry operated on.

We may add up these four flops per entry by geometric reasoning, as in Lecture 8. Each successive step of the outer loop operates on fewer rows, because during step k , rows $1, \dots, k-1$ are not changed. Furthermore, each step operates on fewer columns, because columns $1, \dots, k-1$ of the rows operated on are zero and are skipped. Thus the work done by one outer step can be represented by a single layer of the following solid:



The total number of operations corresponds to four times the volume of the solid. To determine the volume pictorially we may divide the solid into two pieces:



The solid on the left has the shape of a ziggurat and converges to a pyramid as $n \rightarrow \infty$, with volume $\frac{1}{3}n^3$. The solid on the right has the shape of a staircase and converges to a prism as $m, n \rightarrow \infty$, with volume $\frac{1}{2}(m-n)n^2$. Combined, the volume is $\sim \frac{1}{2}mn^2 - \frac{1}{6}n^3$. Multiplying by four flops per unit volume, we find

$$\text{Work for Householder orthogonalization: } \sim 2mn^2 - \frac{2}{3}n^3 \text{ flops.} \quad (10.9)$$

Exercises

10.1. Determine the (a) eigenvalues, (b) determinant, and (c) singular values of a Householder reflector. For the eigenvalues, give a geometric argument as well as an algebraic proof.

10.2. (a) Write a MATLAB function `[W,R] = house(A)` that computes an implicit representation of a full QR factorization $A = QR$ of an $m \times n$ matrix A with $m \geq n$ using Householder reflections. The output variables are a lower-triangular matrix $W \in \mathbb{C}^{m \times n}$ whose columns are the vectors v_k defining the successive Householder reflections, and a triangular matrix $R \in \mathbb{C}^{n,n}$.

(b) Write a MATLAB function `Q = formQ(W)` that takes the matrix W produced by `house` as input and generates a corresponding $m \times m$ orthogonal matrix Q .

10.3. Let Z be the matrix

$$Z = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 7 \\ 4 & 2 & 3 \\ 4 & 2 & 2 \end{bmatrix}.$$

Compute three reduced QR factorizations of Z in MATLAB: by the Gram–Schmidt routine `mgs` of Exercise 8.2, by the Householder routines `house` and `formQ` of Exercise 10.2, and by MATLAB’s built-in command `[Q,R] = qr(Z,0)`. Compare these three and comment on any differences you see.

10.4. Consider the 2×2 orthogonal matrices

$$F = \begin{bmatrix} -c & s \\ s & c \end{bmatrix}, \quad J = \begin{bmatrix} c & s \\ -s & c \end{bmatrix}, \quad (10.10)$$

where $s = \sin \theta$ and $c = \cos \theta$ for some θ . The first matrix has $\det F = -1$ and is a reflector—the special case of a Householder reflector in dimension 2. The second has $\det J = 1$ and effects a rotation instead of a reflection. Such a matrix is called a *Givens rotation*.

(a) Describe exactly what geometric effects left-multiplications by F and J have on the plane $\mathbb{R}^2$. (J rotates the plane by the angle θ , for example, but is the rotation clockwise or counterclockwise?)

(b) Describe an algorithm for QR factorization that is analogous to Algorithm 10.1 but based on Givens rotations instead of Householder reflections.

(c) Show that your algorithm involves six flops per entry operated on rather than four, so that the asymptotic operation count is 50% greater than (10.9).

Lecture 11. Least Squares Problems

Least squares data-fitting has been an indispensable tool since its invention by Gauss and Legendre around 1800, with ramifications extending throughout the mathematical sciences. In the language of linear algebra, the problem here is the solution of an overdetermined system of equations $Ax = b$ —rectangular, with more rows than columns. The least squares idea is to “solve” such a system by minimizing the 2-norm of the residual $b - Ax$.

The Problem

Consider a linear system of equations having n unknowns but $m > n$ equations. Symbolically, we wish to find a vector $x \in \mathbb{C}^n$ that satisfies $Ax = b$, where $A \in \mathbb{C}^{m \times n}$ and $b \in \mathbb{C}^m$. In general, such a problem has no solution. A suitable vector x exists only if b lies in $\text{range}(A)$, and since b is an m -vector, whereas $\text{range}(A)$ is of dimension at most n , this is true only for exceptional choices of b . We say that a rectangular system of equations with $m > n$ is *overdetermined*. The vector known as the *residual*,

$$r = b - Ax \in \mathbb{C}^m, \tag{11.1}$$

can perhaps be made quite small by a suitable choice of x , but in general it cannot be made equal to zero.

What can it mean to solve a problem that has no solution? In the case of an overdetermined system of equations, there is a natural answer to this question. Since the residual r cannot be made to be zero, let us instead make

it as small as possible. Measuring the smallness of r entails choosing a norm. If we choose the 2-norm, the problem takes the following form:

$$\begin{aligned} &\text{Given } A \in \mathbb{C}^{m \times n}, m \geq n, b \in \mathbb{C}^m, \\ &\text{find } x \in \mathbb{C}^n \text{ such that } \|b - Ax\|_2 \text{ is minimized.} \end{aligned} \quad (11.2)$$

This is our formulation of the general (linear) least squares problem. The choice of the 2-norm can be defended by various geometric and statistical arguments, and, as we shall see, it certainly leads to simple algorithms—ultimately because the derivative of a quadratic function, which must be set to zero for minimization, is linear.

The 2-norm corresponds to Euclidean distance, so there is a simple geometric interpretation of (11.2). We seek a vector $x \in \mathbb{C}^n$ such that the vector $Ax \in \mathbb{C}^m$ is the closest point in $\text{range}(A)$ to b .

Example: Polynomial Data-Fitting

As an example, let us compare polynomial interpolation, which leads to a square system of equations, and least squares polynomial data-fitting, where the system is rectangular.

Example 11.1. Polynomial Interpolation. Suppose we are given m distinct points $x_1, \dots, x_m \in \mathbb{C}$ and data $y_1, \dots, y_m \in \mathbb{C}$ at these points. Then there exists a unique *polynomial interpolant* to these data in these points, that is, a polynomial of degree at most $m - 1$,

$$p(x) = c_0 + c_1x + \dots + c_{m-1}x^{m-1}, \quad (11.3)$$

with the property that at each x_i , $p(x_i) = y_i$. The relationship of the data $\{x_i\}$, $\{y_i\}$ to the coefficients $\{c_i\}$ can be expressed by the square Vandermonde system seen already in Example 1.1:

$$\begin{bmatrix} 1 & x_1 & x_1^2 & \dots & x_1^{m-1} \\ 1 & x_2 & x_2^2 & \dots & x_2^{m-1} \\ 1 & x_3 & x_3^2 & \dots & x_3^{m-1} \\ \vdots & & & & \vdots \\ 1 & x_m & x_m^2 & \dots & x_m^{m-1} \end{bmatrix} \begin{bmatrix} c_0 \\ c_1 \\ c_2 \\ \vdots \\ c_{m-1} \end{bmatrix} = \begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ \vdots \\ y_m \end{bmatrix}. \quad (11.4)$$

To determine the coefficients $\{c_i\}$ for a given set of data, we can solve this system of equations, which is guaranteed to be nonsingular as long as the points $\{x_i\}$ are distinct (Exercise 37.3).

Figure 11.1 presents an example of this process of polynomial interpolation. We have eleven data points in the form of a discrete square wave, represented

by crosses, and the curve $p(x)$ passes through them, as it must. However, the fit is not at all pleasing. Near the ends of the interval, $p(x)$ exhibits large oscillations that are clearly an artifact of the interpolation process, not a reasonable reflection of the data.




Figure 11.1. Degree 10 polynomial interpolant to eleven data points. The axis scales are not given, as these have no effect on the picture.

This unsatisfactory behavior is typical of polynomial interpolation. The fits it produces are often bad, and they tend to get worse rather than better if more data are utilized. Even if the fit is good, the interpolation process may be ill-conditioned, i.e., sensitive to perturbations of the data (next lecture). To avoid these problems, one can utilize a nonuniform set of interpolation points such as Chebyshev points in the interval $[-1, 1]$. In applications, however, it will not always be possible to choose the interpolation points at will. $\square$

Example 11.2. Polynomial Least Squares Fitting. Without changing the data points, we can do better by reducing the degree of the polynomial. Given $x_1, \dots, x_m$ and $y_1, \dots, y_m$ again, consider now a degree $n-1$ polynomial

$$p(x) = c_0 + c_1x + \dots + c_{n-1}x^{n-1} \quad (11.5)$$

for some $n < m$. Such a polynomial is a least squares fit to the data if it minimizes the sum of the squares of the deviation from the data,

$$\sum_{i=1}^m |p(x_i) - y_i|^2. \quad (11.6)$$

This sum of squares is equal to the square of the norm of the residual, $\|r\|_2^2$, for the rectangular Vandermonde system

$$\begin{bmatrix} 1 & x_1 & & x_1^{n-1} \\ 1 & x_2 & \cdots & x_2^{n-1} \\ 1 & x_3 & & x_3^{n-1} \\ & \vdots & & \vdots \\ 1 & x_m & \cdots & x_m^{n-1} \end{bmatrix} \begin{bmatrix} c_0 \\ c_1 \\ \vdots \\ c_{n-1} \end{bmatrix} \approx \begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ \vdots \\ y_m \end{bmatrix}. \quad (11.7)$$

Figure 11.2 illustrates what we get if we fit the same eleven data points from the last example with a polynomial of degree 7. The new polynomial does not interpolate the data, but it captures their overall behavior much better than the polynomial of Example 11.1. Though one cannot see this in the figure, it is also less sensitive to perturbations. $\square$




Figure 11.2. Degree 7 polynomial least squares fit to the same eleven data points.

Orthogonal Projection and the Normal Equations

How was Figure 11.2 computed? How are least squares problems solved in general? The key to deriving algorithms is orthogonal projection.

The idea is illustrated in Figure 11.3. Our goal is to find the closest point Ax in $\text{range}(A)$ to b , so that the norm of the residual $r = b - Ax$ is minimized. It is clear geometrically that this will occur provided $Ax = Pb$, where $P \in \mathbb{C}^{m \times m}$ is the orthogonal projector (Lecture 6) that maps $\mathbb{C}^m$ onto $\text{range}(A)$. In other words, *the residual $r = b - Ax$ must be orthogonal to $\text{range}(A)$* . We formulate this condition as the following theorem.

Theorem 11.1. *Let $A \in \mathbb{C}^{m \times n}$ ($m \geq n$) and $b \in \mathbb{C}^m$ be given. A vector $x \in \mathbb{C}^n$ minimizes the residual norm $\|r\|_2 = \|b - Ax\|_2$, thereby solving the least squares problem (11.2), if and only if $r \perp \text{range}(A)$, that is,*

$$A^*r = 0, \quad (11.8)$$




Figure 11.3. Formulation of the least squares problem (11.2) in terms of orthogonal projection.

or equivalently,

$$\underline{A^*Ax = A^*b}, \quad (11.9)$$

or again equivalently,

$$\underline{Pb = Ax}, \quad (11.10)$$

where $P \in \mathbb{C}^{m \times m}$ is the orthogonal projector onto $\text{range}(A)$. The $n \times n$ system of equations (11.9), known as the normal equations, is nonsingular if and only if A has full rank. Consequently the solution x is unique if and only if A has full rank.

Proof. The equivalence of (11.8) and (11.10) follows from the properties of orthogonal projectors discussed in Lecture 6, and the equivalence of (11.8) and (11.9) follows from the definition of r . To show that $y = Pb$ is the unique point in $\text{range}(A)$ that minimizes $\|b - y\|_2$, suppose $z \neq y$ is another point in $\text{range}(A)$. Since $z - y$ is orthogonal to $b - y$, the Pythagorean theorem (Exercise 2.2) gives $\|b - z\|_2^2 = \|b - y\|_2^2 + \|y - z\|_2^2 > \|b - y\|_2^2$, as required. Finally, we note that if A^*A is singular, then $A^*Ax = 0$ for some nonzero x , implying $x^*A^*Ax = 0$ (see Exercise 6.3). Thus $Ax = 0$, which implies that A is rank-deficient. Conversely, if A is rank-deficient, then $Ax = 0$ for some nonzero x , implying $A^*Ax = 0$ also, so A^*A is singular. By (11.9), this characterization of nonsingular matrices A^*A implies the statement about the uniqueness of x . $\square$

Pseudoinverse

We have just seen that if A has full rank, then the solution x to the least squares problem (11.2) is unique and is given by $x = (A^*A)^{-1}A^*b$. The matrix

$(A^*A)^{-1}A^*$ is known as the *pseudoinverse* of A , denoted by A^+ :

$$A^+ = (A^*A)^{-1}A^* \in \mathbb{C}^{n,m}. \quad (11.11)$$

This matrix maps vectors $b \in \mathbb{C}^m$ to vectors $x \in \mathbb{C}^n$, which explains why it has dimensions $n \times m$ —more columns than rows.

We can summarize the full-rank linear least squares problem (11.2) as follows. The problem is to compute one or both of the vectors

$$x = A^+b, \quad y = Pb, \quad (11.12)$$

where A^+ is the pseudoinverse of A and P is the orthogonal projector onto $\text{range}(A)$. We now describe the three leading algorithms for doing this.

Normal Equations

The classical way to solve least squares problems is to solve the normal equations (11.9). If A has full rank, this is a square, hermitian positive definite system of equations of dimension n . The standard method of solving such a system is by *Cholesky factorization*, discussed in Lecture 23. This method constructs a factorization $A^*A = R^*R$, where R is upper-triangular, reducing (11.9) to the equations

$$R^*Rx = A^*b. \quad (11.13)$$

Here is the algorithm.

Algorithm 11.1. Least Squares via Normal Equations

1. Form the matrix A^*A and the vector A^*b .
2. Compute the Cholesky factorization $A^*A = R^*R$.
3. Solve the lower-triangular system $R^*w = A^*b$ for w .
4. Solve the upper-triangular system $Rx = w$ for x .

The steps that dominate the work for this computation are the first two (for steps 3 and 4, see Lecture 17). Because of symmetry, the computation of A^*A requires only mn^2 flops, half what the cost would be if A and A^* were arbitrary matrices of the same dimensions. Cholesky factorization, which also exploits symmetry, requires $n^3/3$ flops. All together, solving least squares problems by the normal equations involves the following total operation count:

$$\text{Work for Algorithm 11.1: } \sim mn^2 + \frac{1}{3}n^3 \text{ flops.} \quad (11.14)$$

QR Factorization

The “modern classical” method for solving least squares problems, popular since the 1960s, is based upon reduced QR factorization. By Gram–Schmidt orthogonalization or, more usually, Householder triangularization, one constructs a factorization $A = \hat{Q}\hat{R}$. The orthogonal projector P can then be written $P = \hat{Q}\hat{Q}^*$ (6.6), so we have

$$y = Pb = \hat{Q}\hat{Q}^*b. \quad (11.15)$$

Since $y \in \text{range}(A)$, the system $Ax = y$ has an exact solution. Combining the QR factorization and (11.15) gives

$$\hat{Q}\hat{R}x = \hat{Q}\hat{Q}^*b, \quad (11.16)$$

and left-multiplication by $\hat{Q}^*$ results in

$$\hat{R}x = \hat{Q}^*b. \quad (11.17)$$

(Multiplying by $\hat{R}^{-1}$ now gives the formula $A^+ = \hat{R}^{-1}\hat{Q}^*$ for the pseudoinverse.) Equation (11.17) is an upper-triangular system, nonsingular if A has full rank, and it is readily solved by back substitution (Lecture 17).

Algorithm 11.2. Least Squares via QR Factorization

1. Compute the reduced QR factorization $A = \hat{Q}\hat{R}$.
2. Compute the vector $\hat{Q}^*b$.
3. Solve the upper-triangular system $\hat{R}x = \hat{Q}^*b$ for x .

Notice that (11.17) can also be derived from the normal equations. If $A^*Ax = A^*b$, then $\hat{R}^*\hat{Q}^*\hat{Q}\hat{R}x = \hat{R}^*\hat{Q}^*b$, which implies $\hat{R}x = \hat{Q}^*b$.

The work for Algorithm 11.2 is dominated by the cost of the QR factorization. If Householder reflections are used for this step, we have from (10.9)

$$\text{Work for Algorithm 11.2: } \sim 2mn^2 - \frac{2}{3}n^3 \text{ flops.} \quad (11.18)$$

SVD

In Lecture 31 we shall describe an algorithm for computing the reduced singular value decomposition $A = \hat{U}\hat{\Sigma}V^*$. This suggests another method for solving least squares problems. Now P is represented in the form $P = \hat{U}\hat{U}^*$, giving

$$y = Pb = \hat{U}\hat{U}^*b, \quad (11.19)$$

and the analogues of (11.16) and (11.17) are

$$\hat{U}\hat{\Sigma}V^*x = \hat{U}\hat{U}^*b \quad (11.20)$$

and

$$\hat{\Sigma}V^*x = \hat{U}^*b. \quad (11.21)$$

(Multiplying by $V\hat{\Sigma}^{-1}$ gives $A^+ = V\hat{\Sigma}^{-1}\hat{U}^*$.) The algorithm looks like this.

Algorithm 11.3. Least Squares via SVD

1. Compute the reduced SVD $A = \hat{U}\hat{\Sigma}V^*$.
2. Compute the vector $\hat{U}^*b$.
3. Solve the diagonal system $\hat{\Sigma}w = \hat{U}^*b$ for w .
4. Set $x = Vw$.

Note that whereas QR factorization reduces the least squares problem to a triangular system of equations, the SVD reduces it to a diagonal system of equations, which is of course trivially solved. If A has full rank, the diagonal system is nonsingular.

As before, (11.21) can be derived from the normal equations. If $A^*Ax = A^*b$, then $V\hat{\Sigma}^*\hat{U}^*\hat{U}\hat{\Sigma}V^*x = V\hat{\Sigma}^*\hat{U}^*b$, implying $\hat{\Sigma}V^*x = \hat{U}^*b$.

The operation count for Algorithm 11.3 is dominated by the computation of the SVD. As we shall see in Lecture 31, for $m \gg n$ this cost is approximately the same as for QR factorization, but for $m \approx n$ the SVD is more expensive. A typical estimate is

$$\text{Work for Algorithm 11.3: } \sim 2mn^2 + 11n^3 \text{ flops,} \quad (11.22)$$

but see Lecture 31 for qualifications of this result.

Comparison of Algorithms

Each of the methods we have described is advantageous in certain situations. When speed is the only consideration, Algorithm 11.1 may be the best. However, solving the normal equations is not always stable in the presence of rounding errors, and thus for many years, numerical analysts have recommended Algorithm 11.2 instead as the standard method for least squares problems. This is indeed a natural and elegant algorithm, and we recommend it for “daily use.” If A is close to rank-deficient, however, it turns out that Algorithm 11.2 itself has less-than-ideal stability properties, and in such cases there are good reasons to turn to Algorithm 11.3, based on the SVD.

What are these stability considerations that make one algorithm better than another in some circumstances yet not in others? It is time now to undertake a systematic discussion of such matters. We shall return to the study of algorithms for least squares problems in Lectures 18 and 19.

Exercises

11.1. Suppose the $m \times n$ matrix A has the form

$$A = \begin{bmatrix} A_1 \\ A_2 \end{bmatrix},$$

where A_1 is a nonsingular matrix of dimension $n \times n$ and A_2 is an arbitrary matrix of dimension $(m - n) \times n$. Prove that $\|A^+\|_2 \leq \|A_1^{-1}\|_2$.

11.2. (a) How closely, as measured in the L^2 norm on the interval $[1, 2]$, can the function $f(x) = x^{-1}$ be fitted by a linear combination of the functions e^x , $\sin x$, and $\Gamma(x)$? ($\Gamma(x)$ is the gamma function, a built-in function in MATLAB.) Write a program that determines the answer to at least two digits of relative accuracy using a discretization of $[1, 2]$ and a discrete least squares problem. Write down your estimate of the answer and also of the coefficients of the optimal linear combination, and produce a plot of the optimal approximation. (b) Now repeat, but with $[1, 2]$ replaced by $[0, 1]$. You may find the following fact helpful: if $g(x) = 1/\Gamma(x)$, then $g'(0) = 1$.

11.3. Take $m = 50$, $n = 12$. Using MATLAB's `linspace`, define t to be the m -vector corresponding to linearly spaced grid points from 0 to 1. Using MATLAB's `vander` and `fliplr`, define A to be the $m \times n$ matrix associated with least squares fitting on this grid by a polynomial of degree $n - 1$. Take b to be the function $\cos(4t)$ evaluated on the grid. Now, calculate and print (to sixteen-digit precision) the least squares coefficient vector x by six methods:

- (a) Formation and solution of the normal equations, using MATLAB's `\`,
- (b) QR factorization computed by `mgs` (modified Gram–Schmidt, Exercise 8.2),
- (c) QR factorization computed by `house` (Householder triangularization, Exercise 10.2),
- (d) QR factorization computed by MATLAB's `qr` (also Householder triangularization),
- (e) `x = A\b` in MATLAB (also based on QR factorization),
- (f) SVD, using MATLAB's `svd`.
- (g) The calculations above will produce six lists of twelve coefficients. In each list, shade with red pen the digits that appear to be wrong (affected by rounding error). Comment on what differences you observe. Do the normal equations exhibit instability? You do not have to explain your observations.